

Towards a D3Q27 Lattice Boltzmann Method Implementation Using CUDA

Erik Brobyn

City College of New York

Department of Physics

First Draft May 5, 2020

Abstract

In this report, the Lattice Boltzmann method is investigated for use with CUDA devices in a Windows 10 environment. A sample application is created, and techniques are explored using NVIDIA's CUDA Samples v10.2 as a guide to best-practices. An existing CUDA implementation of a D2Q9 Lattice Boltzmann is analyzed as a potential starting point for further investigation and future modification into a D3Q27 version of the model, and future changes are discussed.

Introduction

The Lattice Boltzmann methods that are used for computational fluid models are popular and well-suited for numerous situations where atmospheric modeling and analysis is warranted. <<history and more about Lattice Boltzmann...>>

Software Investigated

- Microsoft Visual Studio Community 2017 and 2019
- NVIDIA GPU Computing Toolkit (CUDA) v10.2.89 441.22 for Windows 10
- NVIDIA Corporation's CUDA Samples v10.2
- LATTICE BOLTZMANN SIMULATOR GPU accelerated with CUDA
by Tom Scherlis and Henry Friedlander (2017)

About NVIDIA CUDA Samples v10.2

The CUDA Samples is a collection of 176 projects spread across several subject areas, including Graphics, Finance, Simulation, and Imaging. The sample code is mostly written in C, with some C++ used, and many of the samples are fairly short and digestible. There is often brief commented discussion of best-practices for structure and style.

Given its instructional nature, the CUDA Samples are a pretty good standard for GPU-related development best-practices for this project. There is also a crash course in OpenGL API functionality implicit in the graphical components of many CUDA samples. Generally speaking, the CUDA

graphics sample projects have two pieces, a CPU layer, and a GPU or kernel layer.

The CPU layer handles the windowing functions, like OpenGL Utility Toolkit (GLUT) calls, UI handlers, and other display-related functions such as buffer allocation and initialization and other OpenGL calls. This layer is typically implemented in a standard C++ .cpp implementation file.

On the other hand, the GPU or kernel layer, looks somewhat different. Functions in this layer generally have declaration specifiers, `__device__` or `__global__` prefixes, indicating whether or not the function is a kernel function (`__global__`) or meant to be run on a single GPU core (`__device__`). Calls to kernels require the `'<<<'` and `'>>>'` notation; within these brackets blocks and thread sizes must be provided. These kernel layers are typically defined in CUDA-specific .cu implementation files.

As a preliminary for this project, all the CUDA v10.2 samples were compiled and built using both 2017 and 2019 versions of Visual Studio Community Edition. 173 projects were built without errors, and there were 6 errors spread across 3 projects. Two of these projects related to missing Vulkan libraries, the other error was a heap exception.

CUDA Errors Reported by Intellisense

CUDA projects and files are compiled using NVIDIA's compiler, `nvcc.exe`. Since it is not Visual Studio's standard compiler (`cl.exe`), the Visual Studio IDE's built-in code-highlighting tool (commonly referred to as **Intellisense**) is not configured to gracefully handle the nuances of CUDA syntax.

Both the 2017 and 2019 versions of Intellisense found undeclared identifier errors in almost every CUDA Sample project. These errors appear for many CUDA-specific objects, including `threadIdx`, `blockIdx`, some CUDA functions, '<<<' and '>>>' operators, and some variables and other references. Short of turning off Intellisense altogether, the best you can do to fix these errors is to include the following header files in any file that uses CUDA syntax or objects:

```
#include <cuda.h>
#include <cuda_runtime.h>
#include <device_launch_parameters.h>
```

However, this does not fix all the errors Intellisense finds. The '<<<', '>>>' operators and some other symbols continue to be a nuisance, but again, these errors ultimately have no effect on the build.

About `lbm_cuda`

lbm_cuda is a Visual Studio 2019 solution created with a default CUDA 10.2 project. All projects in the solution are set to `Debug` configuration mode, and most of the conventions used by CUDA Samples for v10.2 were used for the projects contained within. They include:

- additional include directories
 `../common/inc`
- lib files location
 `../common/lib`
- output directories for executables
 `../bin/win64/Debug`
- to avoid run-time errors, two dlls are required in
 `../bin/win64/Debug`
 `glew64.dll`
 `freeglut.dll`

These changes are made in each project's **Properties**. In **Solution Explorer**, right-click the project, select **Properties** and expand

Configuration Properties. Find **Output Directory** in the **General** tab, **Additional Include Directories** in the **C/C++>General** tab, and **Additional Library Directories** in the **Linker>General** tab.

The **fluidsGL** CUDA Samples v10.2 project was added to the project space to provide a comparable example of an OpenGL and CUDA-based fluid simulation. It uses the CUFFT library and was only slightly edited for clarification of some concepts. A GLUT-based window management implementation (in a `.cpp` file) makes calls to the GPU (also called kernel) layer via `extern "C"` function prototypes declared in the `.cpp` file. The `extern "C"` functions with CUDA-specific syntax is then defined in a CUDA-specific `.cu` implementation file. This separation marks a clean logical separation between the CPU and kernel layers.

Then the **lbm** project was created from another default CUDA v10.2 project. A CUDA source file `lbm.cu` was created and was initially taken whole from the Lattice Boltzmann Simulator by Tom Scherlis and Henry Friedlander (2017).

Settings for the `lbm` project were then configured to match conventions in the CUDA Samples v10.2 projects. The library `glew64.lib` was also added to **Configuration Properties>Linker>Input>Additional Dependencies**, although its purpose is unclear and may be removed.

Scherlis and Friedlander's D2Q9 Simulator

The original, unmodified version of Scherlis and Friedlander's code was written as a single `.cu` implementation file 1167 lines long. There are several global variables and definitions shared between the CPU and GPU layers.

For this project, this code was first refactored into multiple files with a kernel-based file structure. In addition to providing some improved organization, these steps were also helpful in getting a better understanding of the underlying structure of the program. Some of these changes include:

- Defines and structs shared between the CPU and GPU were split into a separate header file.

- Global variables were moved into a new .cpp file and references to those variables marked extern in the .cu file.
- GPU entry points were declared as extern function prototypes in the .cpp file, and then defined in the .cu file.
- All CPU-related functions were then moved to the .cpp file.
- Two new GPU-related functions were created for handling initialization and cleanup in the new, more modular file structure.
- Original comments were edited and the code reformatted.
- After clean-up, the extern variables were progressively removed entirely in favor of local parameters, passed to extern functions. A pure class-based C++ implementation gradually revealed itself.

Problems with Texture Mapping

The core of the Scherlis and Friedlander's simulator is a 2D texture mapping via an OpenGL pixel buffer object, which is then used with a CUDA graphics buffer:

```
glGenBuffers(1, &pbo);
glBindBuffer(GL_PIXEL_UNPACK_BUFFER, pbo);
...
cudaGraphicsGLRegisterBuffer(
    &cuda_pbo_resource, pbo,
    cudaGraphicsMapFlagsWriteDiscard);
```

This is not well suited for visualization in three dimensions, and some effort was required to replace the texture mapping component with a vertex-based approach. Two examples were used to implement these changes.

The first is a simple OpenGL v1.1 Programming Guide sample called checker.c. This example creates and maps a checkerboard texture to two GL_QUADS rendered in three-dimensional space. It accomplishes this with calls to glMatrixMode(GL_MODELVIEW) to set up the viewport, and then gluPerspective to initialize the camera's frustum parameters. Adding this functionality to lbm creates a 3D view of the generated field's texture, but it still remains a 2D texture (albeit rendered in 3D).

There are two goals for this project. The first goal is to create an additional dimension visible through the viewport. To this end, various samples were tested to find a useful mouse-controlled 3D viewport component for OpenGL/GLUT, and this functionality was gradually added to the lbm project.

The second goal of this project is to enhance the Scherlis and Friedlander code to include an additional third dimension to the lattice itself, along with new velocities, ultimately yielding a D3Q27 model visualization in three dimensions.

Comparison of fluidsGL and lbm Kernels

fluidsGL is an implementation of the stable fluids FFT-based algorithm by Jos Stam. The core kernel algorithm for fluidsGL is invoked in the simulateFluids function:

```
void simulateFluids(void)
{
    // simulate fluid
    advectVelocity(dvfield, (float *)vxfield, (float *)vyfield,
        DIM, RPAW, DIM, DT);
    diffuseProject(vxfield, vyfield, CPAW, DIM, DT, VIS);
    updateVelocity(dvfield, (float *)vxfield, (float *)vyfield,
        DIM, RPAW, DIM);
    advectParticles(vbo, dvfield, DIM, DIM, DT);
}
```

Here we can identify the key steps of the Lattice Boltzmann kernel update:

- advect velocity
- velocity diffusion and projection
- update velocity
- advect particles

Each of the functions called above have kernel delegates. In lbm, the render function contains its core kernel algorithm. Here is a snippet:

```
lbm_node* before = array1_gpu;
lbm_node* after = array2_gpu;

//determine number of threads and blocks required
dim3 threads_per_block = dim3(32, 32, 1);
dim3 number_of_blocks =
    dim3(params.width / 32 + 1, params.height / 32 + 1, 1);

collide<<<number_of_blocks, threads_per_block>>>
    (d2q9_gpu, before, after, params_gpu, barrier_gpu);

before = array2_gpu;
after = array1_gpu;

stream<<<number_of_blocks, threads_per_block>>>
    (d2q9_gpu, before, after, barrier_gpu, params_gpu);
...

bounceAndRender<<<number_of_blocks, threads_per_block>>>
    (d2q9_gpu, before, after, barrier_gpu, params_gpu,
```

```
d_out, prex, prey);
```

Here we can identify the key steps of the Lattice Boltzmann kernel update:

- collide
- stream
- bounce

Both programs call their kernels in their respective `glutDisplayFunc` handlers.

Results

Conclusions

This section contains the most important of the results, which ought to be interpreted and explained. Any statement you make here has to be well thought in advance and absolutely justified.

List of References

- [1] Nolan Goodnight CUDA/OpenGL Fluid Simulation 2013.
- [2] Stam, J. 1999. "Stable Fluids." In Proceedings of SIGGRAPH 1999.
- [3] checker.c from the OpenGL v1.1 Programming Guide 1997. SGI
<https://www.opengl.org/archives/resources/code/samples/redbook/checker.c>

Error Analysis

Error analysis is beyond the scope of this report.

Appendix A: Code

The code-base is publicly available for review, downloading, or forking at the following address:

https://github.com/charlielobster/lbm_cuda